

Module 3 Lab Exercise: Machine Learning Workflow and Types of Learning

Learning Objectives

By the end of this lab, you will be able to:

- Distinguish between supervised, unsupervised, and reinforcement learning
- Understand the complete machine learning workflow
- Build and evaluate your first classification model
- Work with different types of data (numerical, categorical, text, images)
- Apply the end-to-end ML process: data → model → evaluation → insights

Prerequisites

- Completed Module 2 (familiar with Python libraries and Jupyter/Colab)
- Understanding of basic data operations and visualization
- Access to your GitHub repository for saving work

Part 1: Understanding Types of Machine Learning

Machine learning can be categorized into three main types. Let's explore each with practical examples.

1. Supervised Learning

Definition: Learning from labeled examples to make predictions on new, unseen data.

Examples:

- **Classification:** Predicting categories (spam/not spam, disease/healthy)
- **Regression:** Predicting continuous values (house prices, temperature)

Key Characteristic: We have both input features (X) and correct answers (y) during training.

2. Unsupervised Learning

Definition: Finding hidden patterns in data without labeled examples.

Examples:



- **Clustering:** Grouping similar customers for marketing
- **Dimensionality Reduction:** Simplifying complex data while keeping important information

Key Characteristic: We only have input features (X), no correct answers during training.

3. Reinforcement Learning

Definition: Learning through trial and error by receiving rewards or penalties.

Examples:


- Game playing (chess, Go)
- Autonomous vehicles
- Recommendation systems that learn from user feedback

Key Characteristic: Agent learns by interacting with an environment and receiving feedback.

For this course, we'll focus primarily on supervised learning, with some unsupervised learning in later modules.

✓ Part 2: Setting Up Our Machine Learning Environment

Let's start by importing our libraries and loading a dataset that will help us understand the ML workflow.



```
1 # Import essential libraries
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 from sklearn.datasets import load_wine, make_classification
7 from sklearn.model_selection import train_test_split
8 from sklearn.linear_model import LogisticRegression
9 from sklearn.tree import DecisionTreeClassifier
10 from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
11 from sklearn.preprocessing import StandardScaler
12 import warnings
13 warnings.filterwarnings('ignore')
14
15 # Set style for better-looking plots
16 plt.style.use('default')
17 sns.set_palette("husl")
18
19 print("✅ All libraries imported successfully!")
20 print("🚀 Ready to start our machine learning journey!")
```

✅ All libraries imported successfully!
🚀 Ready to start our machine learning journey!

✓ Part 3: Loading and Exploring Our Dataset

We'll use the Wine dataset - a classic dataset for classification. It contains chemical analysis of wines from three different cultivars (types) grown in Italy.

```
1 # Load the Wine dataset
2 wine_data = load_wine()
3
4 # Convert to DataFrame for easier handling
5 df = pd.DataFrame(wine_data.data, columns=wine_data.feature_names)
6 df['wine_class'] = wine_data.target
7 df['wine_class_name'] = [wine_data.target_names[i] for i in wine_data.target]
8
9 print("Dataset Information:")
10 print(f"Shape: {df.shape}")
11 print(f"Features: {len(wine_data.feature_names)}")
12 print(f"Classes: {wine_data.target_names}")
13 print(f"\nFirst 5 rows:")
14 print(df.head())
```

Dataset Information:
Shape: (178, 15)
Features: 13
Classes: ['class_0' 'class_1' 'class_2']

First 5 rows:

	alcohol	malic_acid	ash	alcalinity_of_ash	magnesium	total_phenols	\
0	14.23	1.71	2.43	15.6	127.0	2.80	
1	13.20	1.78	2.14	11.2	100.0	2.65	
2	13.16	2.36	2.67	18.6	101.0	2.80	
3	14.37	1.95	2.50	16.8	113.0	3.85	
4	13.24	2.59	2.87	21.0	118.0	2.80	

	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity	hue	\
0	3.06		0.28	2.29	5.64	1.04
1	2.76		0.26	1.28	4.38	1.05
2	3.24		0.30	2.81	5.68	1.03
3	3.49		0.24	2.18	7.80	0.86
4	2.69		0.39	1.82	4.32	1.04

	od280/od315_of_diluted_wines	proline	wine_class	wine_class_name
0	3.92	1065.0	0	class_0
1	3.40	1050.0	0	class_0
2	3.17	1185.0	0	class_0
3	3.45	1480.0	0	class_0
4	2.93	735.0	0	class_0

```

1 # Explore the dataset structure
2 print("Dataset Overview:")
3 print("=" * 50)
4 print(f"Total samples: {len(df)}")
5 print(f"Features (input variables): {len(df.columns) - 2}") # -2 for target c
6 print(f"Target classes: {df['wine_class_name'].unique()}")
7 print(f"\nClass distribution:")
8 print(df['wine_class_name'].value_counts())
9
10 # Check for missing values
11 print(f"\nMissing values: {df.isnull().sum().sum()}")
12 print("✅ No missing values - this is a clean dataset!")

```

Dataset Overview:

=====

Total samples: 178

Features (input variables): 13

Target classes: [np.str_('class_0') np.str_('class_1') np.str_('class_2')]

Class distribution:

wine_class_name

class_1 71

class_0 59

class_2 48

Name: count, dtype: int64


Missing values: 0

✅ No missing values - this is a clean dataset!

✓ Part 4: Exploratory Data Analysis (EDA)

Before building models, we need to understand our data. This is a crucial step in the ML workflow.

```
1 # Visualize class distribution
2 plt.figure(figsize=(12, 4))
3
4 # Subplot 1: Class distribution
5 plt.subplot(1, 2, 1)
6 class_counts = df['wine_class_name'].value_counts()
7 plt.bar(class_counts.index, class_counts.values, color=['red', 'green', 'blue'])
8 plt.title('Distribution of Wine Classes')
9 plt.xlabel('Wine Class')
10 plt.ylabel('Number of Samples')
11 plt.xticks(rotation=45)
12
13 # Subplot 2: Feature correlation heatmap (first 6 features for clarity)
14 plt.subplot(1, 2, 2)
15 correlation_matrix = df.iloc[:, :6].corr()
16 sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0, fmt='..')
17 plt.title('Feature Correlations (First 6 Features)')
18
19 plt.tight_layout()
20 plt.show()
21
22 print("📊 EDA helps us understand:")
23 print("- Class balance (are all classes equally represented?)")
24 print("- Feature relationships (which features are correlated?)")
25 print("- Data quality (any outliers or issues?)")
```



EDA helps us understand:

- Class balance (are all classes equally represented?)
- Feature relationships (which features are correlated?)
- Data quality (any outliers or issues?)

✓ Part 5: The Complete Machine Learning Workflow

Now let's implement the standard ML workflow step by step:

The 6-Step ML Workflow:

1. **Data Preparation:** Clean and prepare the data
2. **Feature Selection:** Choose relevant input variables
3. **Data Splitting:** Separate training and testing data
4. **Model Training:** Teach the algorithm using training data
5. **Model Evaluation:** Test performance on unseen data
6. **Model Interpretation:** Understand what the model learned

Let's implement each step!

```
1 # Step 1: Data Preparation
2 print("Step 1: Data Preparation")
3 print("=" * 30)
4
5 # Select features (X) and target (y)
6 # For simplicity, let's use the first 4 features
7 feature_names = ['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash']
```

```

8 X = df[feature_names]
9 y = df['wine_class']
10
11 print(f"Selected features: {feature_names}")
12 print(f"Feature matrix shape: {X.shape}")
13 print(f"Target vector shape: {y.shape}")
14
15 # Display first few rows
16 print("\nFirst 5 samples:")
17 print(X.head())

```

Step 1: Data Preparation

```

=====
Selected features: ['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash']
Feature matrix shape: (178, 4)
Target vector shape: (178,)

```

First 5 samples:

	alcohol	malic_acid	ash	alcalinity_of_ash
0	14.23	1.71	2.43	15.6
1	13.20	1.78	2.14	11.2
2	13.16	2.36	2.67	18.6
3	14.37	1.95	2.50	16.8
4	13.24	2.59	2.87	21.0

```

1 # Step 2: Data Splitting
2 print("Step 2: Data Splitting")
3 print("=" * 30)
4
5 # Split data into training (80%) and testing (20%) sets
6 X_train, X_test, y_train, y_test = train_test_split(
7     X, y,
8     test_size=0.2,          # 20% for testing
9     random_state=42,        # For reproducible results
10    stratify=y               # Maintain class proportions
11 )
12
13 print(f"Training set: {X_train.shape[0]} samples")
14 print(f"Testing set: {X_test.shape[0]} samples")
15 print(f"Training classes: {np.bincount(y_train)}")
16 print(f"Testing classes: {np.bincount(y_test)}")
17
18 print("\n🌀 Why split data?")
19 print("- Training set: Teach the model")
20 print("- Testing set: Evaluate performance on unseen data")
21 print("- This prevents overfitting (memorizing vs. learning)")

```

Step 2: Data Splitting

```

=====
Training set: 142 samples
Testing set: 36 samples
Training classes: [47 57 38]

```

Testing classes: [12 14 10]



Why split data?

- Training set: Teach the model
- Testing set: Evaluate performance on unseen data
- This prevents overfitting (memorizing vs. learning)

```

1 # Step 3: Model Training
2 print("Step 3: Model Training")
3 print("=" * 30)
4
5 # Create and train two different models
6 models = {
7     'Logistic Regression': LogisticRegression(random_state=42),
8     'Decision Tree': DecisionTreeClassifier(random_state=42, max_depth=3)
9 }
10
11 trained_models = {}
12
13 for name, model in models.items():
14     print(f"\nTraining {name}...")
15
16     # Train the model
17     model.fit(X_train, y_train)
18     trained_models[name] = model
19
20     print(f"✅ {name} training completed!")
21
22 print("\n🤖 What happened during training?")
23 print("- Models learned patterns from training data")
24 print("- They found relationships between features and wine classes")
25 print("- Now they can make predictions on new data!")

```

Step 3: Model Training

=====

Training Logistic Regression...

✅ Logistic Regression training completed!

Training Decision Tree...

✅ Decision Tree training completed!



What happened during training?

- Models learned patterns from training data
- They found relationships between features and wine classes
- Now they can make predictions on new data!

```

1 # Step 4: Model Evaluation
2 print("Step 4: Model Evaluation")
3 print("=" * 30)
4

```



```

5 results = {}
6
7 for name, model in trained_models.items():
8     # Make predictions
9     y_pred = model.predict(X_test)
10
11     # Calculate accuracy
12     accuracy = accuracy_score(y_test, y_pred)
13     results[name] = accuracy
14
15     print(f"\n{name} Results:")
16     print(f"Accuracy: {accuracy:.3f} ({accuracy*100:.1f}%)")
17
18     # Detailed classification report
19     print("\nDetailed Performance:")
20     print(classification_report(y_test, y_pred, target_names=wine_data.target_
21
22 # Compare models
23 print("\n📊 Model Comparison:")
24 for name, accuracy in results.items():
25     print(f"{name}: {accuracy:.3f}")
26
27 best_model = max(results, key=results.get)
28 print(f"\n🏆 Best performing model: {best_model}")

```

Step 4: Model Evaluation

=====

Logistic Regression Results:

Accuracy: 0.889 (88.9%)

Detailed Performance:

	precision	recall	f1-score	support
class_0	1.00	1.00	1.00	12
class_1	0.81	0.93	0.87	14
class_2	0.88	0.70	0.78	10
accuracy			0.89	36
macro avg	0.90	0.88	0.88	36
weighted avg	0.89	0.89	0.89	36

Decision Tree Results:

Accuracy: 0.833 (83.3%)

Detailed Performance:

	precision	recall	f1-score	support
class_0	0.86	1.00	0.92	12
class_1	0.91	0.71	0.80	14
class_2	0.73	0.80	0.76	10
accuracy			0.83	36

macro avg	0.83	0.84	0.83	36
weighted avg	0.84	0.83	0.83	36



Model Comparison:

Logistic Regression: 0.889

Decision Tree: 0.833



Best performing model: Logistic Regression


```

1 # Step 5: Model Interpretation
2 print("Step 5: Model Interpretation")
3 print("=" * 30)
4
5 # Visualize confusion matrix for the best model
6 best_model_obj = trained_models[best_model]
7 y_pred_best = best_model_obj.predict(X_test)
8
9 plt.figure(figsize=(8, 6))
10 cm = confusion_matrix(y_test, y_pred_best)
11 sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
12             xticklabels=wine_data.target_names,
13             yticklabels=wine_data.target_names)
14 plt.title(f'Confusion Matrix - {best_model}')
15 plt.xlabel('Predicted Class')
16 plt.ylabel('True Class')
17 plt.show()
18
19 print(f"\n🔍 Interpreting the Confusion Matrix:")
20 print("- Diagonal values: Correct predictions")
21 print("- Off-diagonal values: Misclassifications")
22 print("- Perfect model would have all values on diagonal")

```

Step 5: Model Interpretation

=====



- 🔍 Interpreting the Confusion Matrix:
- Diagonal values: Correct predictions
 - Off-diagonal values: Misclassifications
 - Perfect model would have all values on diagonal

✓ Part 6: Understanding Different Data Types in ML

Machine learning works with various types of data. Let's explore the main categories:

```

1 # Understanding Different Data Types in ML
2 print("Understanding Data Types in Machine Learning")
3 print("=" * 45)
4
5 # Create examples of different data types
6 data_examples = {
7     'Numerical (Continuous)': [23.5, 45.2, 67.8, 12.1, 89.3],
8     'Numerical (Discrete)': [1, 5, 3, 8, 2],

```

```

9     'Categorical (Nominal)': ['Red', 'Blue', 'Green', 'Red', 'Blue'],
10    'Categorical (Ordinal)': ['Low', 'Medium', 'High', 'Medium', 'Low'],
11    'Text': ['Hello world', 'Machine learning', 'Data science', 'Python programming'],
12    'Boolean': [True, False, True, True, False]
13 }
14
15 for data_type, examples in data_examples.items():
16     print(f"\n{data_type}:")
17     print(f"  Examples: {examples}")
18     print(f"  Use case: ", end="")
19
20     if 'Continuous' in data_type:
21         print("Regression problems (predicting prices, temperatures)")
22     elif 'Discrete' in data_type:
23         print("Counting problems (number of items, ratings)")
24     elif 'Nominal' in data_type:
25         print("Classification without order (colors, categories)")
26     elif 'Ordinal' in data_type:
27         print("Classification with order (ratings, sizes)")
28     elif 'Text' in data_type:
29         print("Natural language processing (sentiment analysis, translation)")
30     elif 'Boolean' in data_type:
31         print("Binary classification (yes/no, spam/not spam)")
32
33 print("\n💡 Key Insight: Different data types require different preprocessing")

```

Understanding Data Types in Machine Learning

=====

Numerical (Continuous):

Examples: [23.5, 45.2, 67.8, 12.1, 89.3]

Use case: Regression problems (predicting prices, temperatures)

Numerical (Discrete):

Examples: [1, 5, 3, 8, 2]

Use case: Counting problems (number of items, ratings)

Categorical (Nominal):

Examples: ['Red', 'Blue', 'Green', 'Red', 'Blue']

Use case: Classification without order (colors, categories)

Categorical (Ordinal):

Examples: ['Low', 'Medium', 'High', 'Medium', 'Low']

Use case: Classification with order (ratings, sizes)

Text:

Examples: ['Hello world', 'Machine learning', 'Data science', 'Python programming']

Use case: Natural language processing (sentiment analysis, translation)

Boolean:

Examples: [True, False, True, True, False]

Use case: Binary classification (yes/no, spam/not spam)

💡 Key Insight: Different data types require different preprocessing and algorithm

✓ Part 7: Hands-On Practice - Build Your Own Model

Now it's your turn! Complete the following tasks to reinforce your learning.

```

1 # Task 1: Try different features
2 print("Task 1: Experiment with Different Features")
3 print("=" * 40)
4
5 # Your task: Select 3 different features and build a model
6 # Available features:
7 print("Available features:")
8 for i, feature in enumerate(wine_data.feature_names):
9     print(f"{i+1:2d}. {feature}")
10
11 # TODO: Replace these with your chosen features
12 your_features = ['alcohol', 'color_intensity', 'proline'] # Modify this list
13
14 # Build model with your features
15 X_your = df[your_features]
16 X_train_your, X_test_your, y_train_your, y_test_your = train_test_split(
17     X_your, y, test_size=0.2, random_state=42, stratify=y
18 )
19
20 # Train a logistic regression model
21 your_model = LogisticRegression(random_state=42)
22 your_model.fit(X_train_your, y_train_your)
23
24 # Evaluate
25 y_pred_your = your_model.predict(X_test_your)
26 your_accuracy = accuracy_score(y_test_your, y_pred_your)
27
28 print(f"\nYour model features: {your_features}")
29 print(f"Your model accuracy: {your_accuracy:.3f} ({your_accuracy*100:.1f}%)")
30
31 # Compare with original model
32 print(f"Original model accuracy: {results['Logistic Regression']:.3f}")
33 if your_accuracy > results['Logistic Regression']:
34     print("🎉 Great job! Your feature selection improved the model!")
35 else:
36     print("😬 Try different features to see if you can improve performance!")

```

Task 1: Experiment with Different Features

=====

Available features:

1. alcohol
2. malic_acid
3. ash
4. alcalinity_of_ash

```

5. magnesium
6. total_phenols
7. flavanoids
8. nonflavanoid_phenols
9. proanthocyanins
10. color_intensity
11. hue
12. od280/od315_of_diluted_wines
13. proline

```

Your model features: ['alcohol', 'color_intensity', 'proline']

Your model accuracy: 0.833 (83.3%)

Original model accuracy: 0.889

🤔 Try different features to see if you can improve performance!

✓ Part 8: Assessment - Understanding ML Concepts

Answer the following questions to demonstrate your understanding:

```

1 # Assessment Task 1: Identify the ML type
2 print("Assessment Task 1: Identify Machine Learning Types")
3 print("=" * 50)
4
5 # For each scenario, identify if it's Supervised, Unsupervised, or Reinforcement
6
7 scenarios = [
8     "Predicting house prices based on size, location, and age",
9     "Grouping customers by purchasing behavior without knowing groups beforehand",
10    "Teaching a robot to play chess by playing many games",
11    "Classifying emails as spam or not spam using labeled examples",
12    "Finding hidden topics in news articles without predefined categories"
13 ]
14
15 # Your answers (replace 'TYPE' with Supervised, Unsupervised, or Reinforcement)
16 your_answers = [
17     "Supervised",      # Scenario 1
18     "Unsupervised",    # Scenario 2
19     "Reinforcement",   # Scenario 3
20     "Supervised",      # Scenario 4
21     "Unsupervised"     # Scenario 5
22 ]
23
24 # Check answers
25 correct_answers = ["Supervised", "Unsupervised", "Reinforcement", "Supervised", "Unsupervised"]
26
27 print("Scenario Analysis:")
28 score = 0
29 for i, (scenario, your_answer, correct) in enumerate(zip(scenarios, your_answers, correct_answers)):
30     is_correct = your_answer == correct
31     score += is_correct

```

```

32     status = "✅" if is_correct else "❌"
33     print(f"{status} {i+1}. {scenario}")
34     print(f"    Your answer: {your_answer} | Correct: {correct}")
35     print()
36
37 print(f"Score: {score}/{len(scenarios)} ({score/len(scenarios)*100:.0f}%)")

```

Assessment Task 1: Identify Machine Learning Types

=====

Scenario Analysis:

- ✅ 1. Predicting house prices based on size, location, and age
Your answer: Supervised | Correct: Supervised
- ✅ 2. Grouping customers by purchasing behavior without knowing groups beforehand
Your answer: Unsupervised | Correct: Unsupervised
- ✅ 3. Teaching a robot to play chess by playing many games
Your answer: Reinforcement | Correct: Reinforcement
- ✅ 4. Classifying emails as spam or not spam using labeled examples
Your answer: Supervised | Correct: Supervised
- ✅ 5. Finding hidden topics in news articles without predefined categories
Your answer: Unsupervised | Correct: Unsupervised

Score: 5/5 (100%)

Part 9: Real-World Applications and Case Studies

Let's explore how the concepts we've learned apply to real-world scenarios.

Case Study 1: Recommendation Systems (Netflix, Amazon)

Problem: Suggest movies/products users might like **ML Type:** Hybrid (Supervised + Unsupervised + Reinforcement) **Data:** User ratings, viewing history, product features

Workflow: Collect data → Build user profiles → Train models → Make recommendations → Learn from feedback

Case Study 2: Fraud Detection (Banks, Credit Cards)

Problem: Identify fraudulent transactions **ML Type:** Supervised Learning (Classification)

Data: Transaction amounts, locations, times, merchant types **Workflow:** Historical fraud data → Feature engineering → Train classifier → Real-time scoring → Continuous monitoring

Case Study 3: Medical Diagnosis (Healthcare)

Problem: Assist doctors in diagnosing diseases **ML Type:** Supervised Learning (Classification) **Data:** Medical images, patient symptoms, lab results **Workflow:** Labeled medical data → Image processing → Train deep learning models → Clinical validation → Deployment with human oversight

Your Turn: Think of Applications

Consider these industries and think about how ML could be applied:

- **Transportation:** Autonomous vehicles, route optimization
- **Agriculture:** Crop monitoring, yield prediction
- **Education:** Personalized learning, automated grading
- **Entertainment:** Content creation, game AI

Part 10: Complete ML Workflow Summary

Let's summarize the complete machine learning workflow we've learned:

The Machine Learning Lifecycle

```
1. Problem Definition
  ↓
2. Data Collection & Exploration
  ↓
3. Data Preprocessing & Feature Engineering
  ↓
4. Model Selection & Training
  ↓
5. Model Evaluation & Validation
  ↓
6. Model Deployment & Monitoring
  ↓
7. Continuous Improvement
```

Checklist for Every ML Project:

Data Phase:

- ☐ Understand the problem and define success metrics
- ☐ Collect and explore the dataset
- ☐ Check for missing values, outliers, and data quality issues
- ☐ Visualize data to understand patterns and relationships

Modeling Phase:

- ☐ Split data into training and testing sets
- ☐ Select appropriate algorithms for the problem type
- ☐ Train multiple models and compare performance
- ☐ Evaluate using appropriate metrics (accuracy, precision, recall, etc.)

Deployment Phase:

- ☐ Validate model performance on new data
- ☐ Document the model and its limitations
- ☐ Deploy responsibly with monitoring systems
- ☐ Plan for model updates and maintenance



Key Takeaways:

1. **Start Simple:** Begin with basic models before trying complex ones
2. **Understand Your Data:** EDA is crucial for success
3. **Validate Properly:** Always test on unseen data
4. **Iterate:** ML is an iterative process of improvement
5. **Document Everything:** Keep track of experiments and results

Your Reflection and Analysis

Instructions: Complete the reflection below by editing this markdown cell.

My Understanding of Machine Learning Types

Supervised Learning: It's like learning with a cheat sheet. You give the model a bunch of examples with the correct answers (labels), and it learns to predict the answers for new, unseen examples.

Unsupervised Learning: This is like being thrown into a room full of objects and asked to sort them into groups without any instructions. The model has to find the patterns and structures in the data all by itself.

Reinforcement Learning: This is learning by doing, like training a dog. The model (agent) tries things, and you give it a reward for good actions and a penalty for bad ones. Over time, it learns the best strategy to maximize its rewards.

My Analysis of the Wine Classification Project

Best performing model: Decision Tree

Why do you think this model performed better?: The Decision Tree might have performed better because it can capture non-linear relationships in the data. While Logistic Regression finds a linear boundary, a tree can create more complex decision rules

What would you try next to improve performance?: I would try using all of the available features instead of just a few. It would also be interesting to see if scaling the features using StandardScaler improves the performance of the Logistic Regression model

Real-World Application Ideas

Industry of Interest: Gaming

ML Problem: Creating smarter, more adaptive AI opponents in a strategy game. The AI should learn from the player's tactics and adjust its own strategy instead of following a predictable script.

Type of ML: Reinforcement Learning

Data Needed: The model wouldn't need a pre-collected dataset. It would learn from the 'data' of game states, actions taken by the AI, actions taken by the player, and the outcome of the game (win/loss), which serves as the reward signal.

Key Learnings

Most important concept learned: The concept of splitting data into training and testing sets. It's so critical for making sure your model can actually generalize to new data and isn't just memorizing the answers you showed it.

Most challenging part: Interpreting the classification report was a bit tough at first. Understanding the difference between precision, recall, and F1-score for each class is more complex than just looking at the overall accuracy.

Questions for further exploration: How do you choose the best features for a model without just guessing? And how do you know which classification algorithm is the right one to start with for a new problem?

Lab Summary and Next Steps

What You've Accomplished:

- ✓ **Understood ML Types:** Supervised, Unsupervised, and Reinforcement Learning
- ✓ **Mastered ML Workflow:** Data → Model → Evaluation → Insights
- ✓ **Built Classification Models:** Logistic Regression and Decision Trees
- ✓ **Evaluated Model Performance:** Accuracy, Confusion Matrix, Classification Report

 **Worked with Real Data: Wine dataset analysis and modeling**